TrustLedger SRM Intelligence Architecture Review — Part 1

Core Intelligence Model

Working Review Version: 0.1 | **Date:** September 2026

Review status

This review is architectural and strategic, not an implementation instruction. It establishes the conceptual foundation against which future TrustLedger capabilities will be evaluated. No development packet should be opened from this document until the architecture review is completed and explicitly approved.

1. Purpose of Part 1

TrustLedger has moved beyond being simply a digital register for stakeholders, engagements, grievances and commitments. The platform contains intelligence-oriented capabilities including stakeholder intelligence, trust intelligence, relationship signals, grievance and case intelligence, ESG/social-performance indicators, engagement planning, evidence-linked reporting, rule-based recommendations and AI-assisted suggestions.

The review therefore asks: **What is the underlying intelligence architecture that allows TrustLedger to continuously develop new intelligence capabilities without repeatedly redesigning the platform?**

2. The Architectural Proposition

TrustLedger should be an evidence-led Stakeholder Relationship Intelligence platform that transforms information about people, relationships, engagements, places, issues and commitments into explainable intelligence that supports better human decisions and actions.

TrustLedger should not attempt to become a generic CRM, GIS platform, ERP, project-management replacement, generic AI platform, complete ESG platform, community-development management system, or repository for every practice associated with SRM.

TrustLedger should occupy the relationship and social intelligence layer around complex projects and programmes.

3. The TrustLedger Operational Spine

Organisation → Project / Programme → Place / Context → Stakeholder → Relationship → Engagement → Commitment → Evidence → Decision → Action → Outcome → Learning.

This is an architectural relationship model. Intelligence must remain anchored to the things that actually happen in stakeholder relationships.

4. Why the Operational Spine Matters

TrustLedger intelligence should remain traceable to operational reality. A statement such as “community sentiment is deteriorating” becomes more useful when the platform can establish which stakeholders, place, project, issue, engagement, evidence, commitments, relationship signals and actions are involved. The distinction is between information and traceable intelligence.

5. The Intelligence Model

DATA → CONTEXT → SIGNAL → INTERPRETATION → INTELLIGENCE → RECOMMENDATION → HUMAN DECISION → ACTION → OUTCOME → LEARNING

This is the proposed TrustLedger Intelligence Cycle.

6. Data

Data represents what TrustLedger knows: stakeholder identity, role, location, project, meeting, engagement, issue, grievance, commitment, attendance, evidence, response, status, dates and documents.

Data answers: What do we know?

7. Context

Context answers: What surrounds the data?

Stakeholder context may include role, authority, influence, interests and organisational or community position. Operating context may include BAU pressure, workload, capacity, competing priorities, availability and accountability. Project, social, political and institutional context may also be relevant.

Stakeholder complexity should therefore be understood as contextual and relational rather than as a property of an isolated stakeholder.

8. Signal

A signal is a meaningful change, pattern or condition detected in data and context.

Examples include declining communication with a high-influence stakeholder; several commitments approaching due dates; high BAU pressure; declining participation; recurring concerns; declining trust; or changes in influence pathways.

The signal identifies something deserving attention; it does not automatically establish its meaning.

9. Interpretation

Interpretation explains a signal while preserving alternative explanations.

Example: three overdue commitments could reflect resistance, insufficient capacity, competing priorities, ambiguous ownership, unrealistic commitments, external dependency or communication failure.

A signal should not automatically become a conclusion.

10. Intelligence

Intelligence is the structured interpretation of evidence and context that becomes useful for decision-making.

Example: *Engagement feasibility is declining because a high-influence stakeholder has high operational workload, multiple competing commitments and limited availability.*

11. Recommendation

TrustLedger may produce a recommended response, such as rescheduling an engagement, reducing information burden or confirming priority between competing commitments.

Recommendations should carry reason, evidence, context, confidence and traceability to the relevant rule, model or methodology.

12. Human Decision

TrustLedger may recommend; it should not silently decide. The human remains responsible for accepting, rejecting, modifying, deferring or escalating a recommendation.

This boundary is particularly important in SRM because stakeholder relationships involve judgement, cultural context, political sensitivity, ethics, power and incomplete information.

13. Action, Outcome and Learning

Actions should be traceable to decisions. Outcomes should be distinguished from activities. A meeting being held is an activity; a concern being reduced or a commitment being resolved is an outcome.

Learning closes the loop: **Evidence** → **Intelligence** → **Decision** → **Action** → **Outcome** → **Learning** → **improved future intelligence.**

This gives TrustLedger the potential to become a system of organisational learning around stakeholder relationships.

14. Proposed Intelligence Domains

Stakeholder Intelligence: who stakeholders are, their roles, influence, interests, operating context and changes over time.

Relationship Intelligence: relationships, influence pathways, trust, tension, collaboration, dependency and network dynamics.

Community Intelligence: community context, local structures, social dynamics, concerns, expectations, informal authority, participation conditions and local knowledge.

Engagement & Participation Intelligence: participation, engagement frequency, quality, burden, willingness, accessibility and effectiveness.

Risk & Early-Warning Intelligence: emerging concerns, deteriorating relationships, unresolved grievances, commitment failure, stakeholder instability and capacity constraints.

Impact & Social Performance Intelligence: social outcomes, distribution of impacts, community benefits, social performance, socioeconomic conditions and mitigation effectiveness.

These are **architectural intelligence domains**, not automatically six new product modules.

15. Decision Intelligence

Decision Intelligence should sit above the domains and ask: **Given everything TrustLedger currently knows, what decision requires attention?**

The long-term opportunity is to move from separate dashboards toward a system that can explain what requires attention, why it matters and what evidence supports the conclusion.

16. Stakeholder Operating Context

Stakeholder understanding should progressively incorporate **Influence, Impact, Capacity and Commitment**, together with BAU pressure, availability, competing priorities, engagement burden and accountability clarity.

These should be contextual attributes that can change over time rather than static scores.

Observed behaviour is not the same as underlying cause. Missing meetings, delayed commitments or low participation may reflect resistance, capacity constraints, competing priorities, poor engagement design, lack of authority or low perceived relevance.

17. AI's Position in the Architecture

AI should not be the intelligence architecture. The architecture should be **Data** → **Context** → **Rules / Models / Methodologies** → **Intelligence** → **AI assistance** → **Human decision**.

AI is an enabler of intelligence, not the source of truth. This protects TrustLedger from dependence on any single AI provider or model.

18. Methodologies Must Be Pluggable

TrustLedger should not hard-code one professional theory of stakeholder management.

Future methodologies may include stakeholder salience, power-interest analysis, influence-network analysis, participation assessment, social-risk analysis, grievance analysis, social performance, community capacity, conflict analysis, social licence, indigenous/local knowledge and client-specific methodologies.

These should operate against the TrustLedger core rather than redefine it.

19. External Data

External data may include demographic, socioeconomic, GIS, environmental, project, survey, public-record, regulatory, client-system or other data.

External data should feed the intelligence architecture rather than redefine the TrustLedger core. Emerging capabilities such as social-media intelligence should therefore be accommodated architecturally without automatically becoming core modules.

20. Layered Intelligence Architecture

The proposed model places **Decision Intelligence** at the top; **Intelligence Domains** beneath it; then **Signal & Interpretation**; **Context**; **Evidence**; and the **Operational Spine** at the foundation.

Across these layers sit methodologies, rules, models and AI, with human governance connecting intelligence to action and outcomes.

21. What This Architecture Deliberately Prevents

Feature accumulation; score accumulation; AI-first architecture; methodology lock-in; dashboard fragmentation; context blindness; and treating activity as impact.

The architecture should keep TrustLedger selective rather than allowing it to become an “everything SRM” system.

22. Key Architectural Principle

TrustLedger Intelligence must be evidence-linked, context-aware, explainable, human-governed and evolution-capable.

Evidence-linked: Where did this come from? **Context-aware:** What circumstances surround it? **Explainable:** Why is TrustLedger saying this? **Human-governed:** Who makes the decision? **Evolution-capable:** Can new professional capabilities be introduced without redesigning the core?

23. What We Should NOT Build Yet

This review does not authorise a new intelligence dashboard, new intelligence DocTypes, a BAU module, predictive AI, stakeholder network visualisation, social-media monitoring, new scoring systems, new AI models, new database architecture, replacement of the Trust layer, or replacement of existing stakeholder CRM capabilities.

These remain questions for later architectural decisions.

24. Preliminary Architectural Principles

1. Core Stability — the operational SRM spine remains stable as capabilities evolve.
2. Intelligence Modularity — intelligence capabilities should be independently extensible.
3. Evidence First — intelligence must be traceable to evidence and observations.
4. Context Matters — stakeholder behaviour must be interpreted within its operating environment.
5. Methodology Agnosticism — methodologies operate on top of the core.
6. Human Governance — AI and analytical systems may recommend; humans remain accountable.
7. No False Certainty — signals, interpretations, hypotheses and verified facts remain distinguishable.
8. Outcome Orientation — measure what changed, not merely what activities occurred.
9. Learning — outcomes should improve future intelligence.
10. Architectural Elasticity — new validated capabilities should be addable without redesigning implemented capabilities.

25. Preliminary TrustLedger Intelligence Cycle

EVIDENCE → CONTEXT → SIGNAL → INTELLIGENCE → RECOMMENDATION → HUMAN DECISION → ACTION → OUTCOME → LEARNING → Context / Intelligence

This feedback loop is the conceptual heart of the architecture.

26. Part 1 Review Conclusion

TrustLedger should not be defined as a collection of SRM features. It should be defined as an **evidence-led relationship intelligence system** built on a stable operational spine.

The architecture should connect evidence, context, relationships, signals, professional methodologies, intelligence, human decisions, actions and outcomes without forcing every emerging SRM practice into the core product.

TrustLedger should not attempt to predict everything about stakeholders. It should build the infrastructure through which evidence, context, relationships, signals, professional methodologies, intelligence, human decisions, actions and outcomes can be connected and progressively learned from.

This provides the foundation for **Part 2: TrustLedger SRM Intelligence Capability Map**.

Appendix — Core Intelligence Cycle

```
EVIDENCE
↓
CONTEXT
↓
SIGNAL
↓
INTERPRETATION
↓
INTELLIGENCE
↓
RECOMMENDATION
↓
HUMAN DECISION
↓
ACTION
↓
OUTCOME
↓
LEARNING
■ feeds future Context / Intelligence
```

Architectural boundary: This document is strategic reference guidance. It does not authorise implementation.