

TrustLedger SRM Intelligence Architecture Review — Part 2

TrustLedger SRM Intelligence Capability Map

Working Review Version: 0.1 | Date: September 2026

Review status

Part 2 translates the Core Intelligence Model established in Part 1 into a capability classification framework. Its purpose is not to create a feature backlog. It provides a disciplined way to decide what belongs inside TrustLedger, what should remain modular or methodological, what should be integrated from elsewhere, and what should deliberately not be built.

1. The Central Question

The SRM profession is broad, multidisciplinary and continuously developing. TrustLedger therefore needs enough architectural breadth to accommodate validated future practices, but enough product discipline to avoid becoming an undifferentiated collection of SRM features.

The objective is architectural comprehensiveness without product indiscrimination.

2. Capability Classification Framework

CORE — Fundamental to TrustLedger's stable operational spine and identity.

MODULE — A substantial user-facing capability that can evolve independently around the core.

INTELLIGENCE — A capability that interprets evidence, context or relationships and produces explainable insight.

METHODOLOGY — A professional framework, analytical method, assessment approach or model that operates on TrustLedger data.

INTEGRATION — External capability or data source that TrustLedger connects to rather than recreates.

CONFIGURATION — Client, project, sector or methodology-specific variation that should be configurable rather than hard-coded.

EXTERNAL — Valuable to SRM practice but better provided by another specialist system or service.

DO NOT BUILD — Outside TrustLedger's strategic role, even if adjacent to SRM.

3. Capability Admission Test

Every proposed capability should pass seven questions:

1. Is it fundamentally related to stakeholder relationships, social performance, participation, community context or trust in complex projects/programmes?
2. Does it strengthen TrustLedger's purpose of building and managing trust?
3. Does it connect naturally to the operational spine?
4. Can it be added without destabilising existing capabilities?
5. Is it better implemented as core, module, intelligence, methodology, integration, configuration or external service?
6. Can its outputs be traced to evidence and context?
7. Is there sufficient professional and user evidence to justify productisation?

4. Core TrustLedger Capability Domain

Organisation & Tenancy — organisation identity, users, roles, permissions, plan/entitlement boundaries and tenant isolation.

Project / Programme — project identity, project context, lifecycle, relationships to places and stakeholder activity.

Place / Context — geographic and contextual reference needed to interpret stakeholder and community information.

Stakeholder — stakeholder identity, roles, affiliations, interests, influence and relevant operating context.

Relationship — relationship state, trust, influence, dependency, collaboration and tension where supported by evidence.

Engagement — interactions, participation, communication, meetings, consultations and engagement outcomes.

Commitment — promises, obligations, ownership, due dates, status and resolution.

Issue / Grievance — concerns, cases, escalation, response, evidence, resolution and learning.

Evidence — documents, observations, records and other evidence linked to operational events.

Decision / Action / Outcome — the human governance layer connecting intelligence to action and measurable outcomes.

Audit / Governance — traceability, permissions, accountability and controlled change.

5. Core Principle: The Core Should Be Small

The core should contain the entities and relationships that remain valid even when professional methods, analytical models or user interfaces change.

For example, **Stakeholder** belongs in the core. A particular stakeholder-prioritisation method does not. Likewise, **Engagement** belongs in the core, while a particular participation-quality framework should be implemented above it.

6. Intelligence Capability Family

Stakeholder Intelligence — stakeholder profiles, influence, interests, roles, capacity, operating context, change over time and prioritisation signals.

Relationship Intelligence — trust, tension, collaboration, influence pathways, dependency and relationship-health signals.

Community Intelligence — community context, local structures, expectations, local knowledge, informal power and social dynamics.

Engagement Intelligence — participation, accessibility, engagement burden, reach, responsiveness, quality and emerging patterns.

Risk & Early Warning Intelligence — emerging stakeholder, relationship, grievance, commitment, participation and capacity signals.

Impact & Social Performance Intelligence — social outcomes, impacts, mitigation, benefits, distribution and performance indicators.

Decision Intelligence — cross-domain prioritisation of what requires human attention and why.

7. Stakeholder Intelligence — Classification

Core: stakeholder identity and fundamental attributes.

Intelligence: influence, relevance, capacity, operating context, changing engagement condition and prioritisation signals.

Methodology: stakeholder salience, power-interest, influence/impact matrices and other professional prioritisation methods.

Configuration: client-defined stakeholder categories, roles, thresholds and assessment dimensions.

Integration: external directories, client CRM or organisational data where appropriate.

Architectural rule: do not make one stakeholder analysis method the permanent TrustLedger data model.

8. Relationship Intelligence — Classification

Core: the existence and traceability of relationships between stakeholders, projects, engagements and relevant records.

Intelligence: relationship health, trust signals, tension, dependency, collaboration and influence pathways.

Methodology: network analysis, relationship assessment, trust models, conflict frameworks and social-licence approaches.

External: specialist network-analysis or graph platforms where complexity exceeds TrustLedger's native role.

Important distinction: a relationship score is not the relationship itself.

9. Community Intelligence — Classification

Core/context: place and community context needed to interpret stakeholder activity.

Intelligence: community dynamics, local concerns, informal authority, participation conditions, community capacity and emerging social patterns.

Methodology: PRA, PLA, CBPR, community profiling, local knowledge and other participatory approaches.

Integration: socioeconomic, demographic, GIS and public datasets.

External: specialist social research or community-development systems when required.

10. Engagement & Participation Intelligence — Classification

Core: engagement events, participants, subjects, actions and outcomes.

Intelligence: participation patterns, engagement quality, burden, accessibility, responsiveness and effectiveness.

Methodology: participation ladders, engagement-quality frameworks, deliberative methods, dialogue approaches and client-specific participation models.

Configuration: engagement types, participation categories, thresholds and assessment criteria.

The platform should distinguish **activity volume** from **participation quality**.

11. Risk & Early-Warning Intelligence — Classification

Core: issues, grievances, commitments, engagement records and evidence from which signals arise.

Intelligence: emerging risk signals, deterioration patterns, unresolved clusters, stakeholder instability, capacity warnings and escalation indicators.

Methodology: social-risk frameworks, escalation models, risk matrices and sector-specific approaches.

Configuration: thresholds, alert rules, severity definitions and escalation policies.

AI: assistance in pattern detection or drafting, never an ungoverned source of truth.

12. Impact & Social Performance Intelligence — Classification

Core: evidence and outcomes associated with project/community interactions.

Intelligence: social performance, impact trends, mitigation effectiveness, benefit distribution and outcome signals.

Methodology: ESG/social-performance frameworks, MEL approaches, impact-assessment frameworks and client-specific indicator systems.

Integration: external socioeconomic, environmental or project datasets.

Configuration: indicator definitions, baselines, targets, reporting periods and materiality criteria.

13. Decision Intelligence — The Coordinating Layer

Decision Intelligence should not become another isolated feature area. It should coordinate signals from the other intelligence domains and answer three questions:

What requires attention?

Why does it matter?

What are the plausible response options?

It should preserve uncertainty and alternative explanations and present evidence so the responsible person can make the decision.

14. Methodology Layer

Methodologies should be treated as reusable analytical lenses over the stable TrustLedger core.

A methodology should define, where applicable: purpose; required inputs; assessment dimensions; calculation or interpretation logic; evidence requirements; confidence/uncertainty; outputs; review cadence; and governance requirements.

This permits TrustLedger to support multiple professional approaches without forcing all users into one theory of SRM.

15. Configuration Layer

Configuration should absorb legitimate variation without code forks. Examples include stakeholder categories, engagement types, issue categories, risk thresholds, indicator sets, reporting periods, assessment scales, project-specific terminology and client-defined workflows.

Rule: configuration should change behaviour within an established capability; it should not be used to disguise a fundamentally new capability.

16. Integration Layer

TrustLedger should integrate with systems that are authoritative or specialised elsewhere: GIS, ERP, CRM, document management, identity, survey systems, socioeconomic data sources, specialist analytics and external AI services where appropriate.

The integration principle is: **connect rather than recreate** when the external system is genuinely better positioned to own the capability.

17. External / Do Not Build Boundary

Potentially external: full ERP, accounting, payroll, enterprise GIS, specialist graph analytics, advanced survey platforms, generic document-management platforms, specialist statistical packages and sector-specific systems.

Do not build simply because SRM professionals use it. A capability belongs in TrustLedger only when it strengthens the platform's strategic role and connects meaningfully to the relationship-intelligence spine.

18. Capability Map — Summary

The following table provides the preliminary classification. It is intentionally directional and should be refined through subsequent architecture review and professional validation.

Capability	Primary class	Secondary treatment	Boundary / rationale
Organisation, tenancy, identity & governance	CORE	Configuration	Stable platform foundation; must remain secure and tenant-aware.
Projects / programmes	CORE	Configuration	Anchors stakeholder and social-performance activity to delivery context.
Place / geographic context	CORE / CONTEXT	Integration	Provides location/context; specialist GIS remains external.
Stakeholder registry	CORE	Intelligence	System of record for stakeholder relationships and related intelligence.
Stakeholder prioritisation	INTELLIGENCE	Methodology	Do not hard-code one prioritisation theory.
Relationship management	CORE + INTELLIGENCE	Methodology	Relationships are core; analytical interpretation is extensible.
Engagement management	CORE / MODULE	Intelligence	Operational record plus participation intelligence.
Commitments	CORE	Intelligence	A key bridge between engagement, accountability and outcomes.
Issues / grievances / cases	CORE / MODULE	Risk intelligence	Operational case handling plus emerging-risk signals.

Capability	Primary class	Secondary treatment	Boundary / rationale
Evidence	CORE	Integration	Traceability layer; external repositories may remain authoritative.
Trust intelligence	INTELLIGENCE	Methodology / rules	Evidence-linked relationship and stability signals.
Community intelligence	INTELLIGENCE	Methodology / Integration	Context and social dynamics; avoid becoming a generic community-development system.
Engagement / participation intelligence	INTELLIGENCE	Methodology / Configuration	Assess quality and feasibility, not just event counts.
Risk & early warning	INTELLIGENCE	Rules / Models / AI assistance	Signals and recommendations; avoid unsupported prediction.
Impact & social performance	INTELLIGENCE / MODULE	Methodology / Integration	Supports social performance and impact without claiming to be all of ESG.
Decision intelligence	INTELLIGENCE	Human governance	Cross-domain coordination of attention and response options.
AI assistance	ENABLER	Integration / Methodology	AI assists interpretation and drafting; it is not the source of truth.
Professional methods / frameworks	METHODOLOGY	Configuration	Pluggable analytical lenses over the stable core.
Client-specific workflows	CONFIGURATION	Module extension	Configure where possible; avoid client-specific forks.
GIS / ERP / specialist systems	INTEGRATION / EXTERNAL	API / data exchange	Connect to authoritative or specialist systems rather than recreate them.
Generic office / productivity suite	EXTERNAL	Integration	Not strategically differentiated by TrustLedger.
Generic accounting / payroll / ERP	DO NOT BUILD	Integration	Adjacent operational domains; outside TrustLedger's relationship-intelligence role.
Generic social-media platform	EXTERNAL	Integration	May provide data/signals; TrustLedger should not become a social network.

19. Capability Independence Test

A new capability should be capable of evolving without forcing redesign of unrelated capabilities.

For example, adding a new community-risk methodology should not require redesigning stakeholder records, engagement records or commitments. Likewise, adding a new social-performance framework should not require replacing the grievance model.

This is the practical meaning of architectural elasticity.

20. The Stable Core / Evolving Edge

Stable core: identity, tenancy, organisation, project/programme, place/context, stakeholder, relationship, engagement, commitment, issue/grievance, evidence, decision, action, outcome, audit and permissions.

Evolving edge: methodologies, intelligence models, scoring approaches, indicators, AI, professional frameworks, reports, analytical tools and external integrations.

The edge can grow rapidly while the core changes slowly and deliberately.

21. Capability Lifecycle

Professional Practice → New Concept → Architectural Review → Capability Classification → Evidence / Validation → Build / Integrate / Configure / Defer / Do Not Build → Implement → Verify → Measure → Learn → Reassess.

This prevents novelty from becoming architecture simply because a new technology or professional trend becomes available.

22. What This Means for Future TrustLedger Development

Future development packets should identify the architectural class of every material capability before implementation begins.

A proposed feature should state: **what problem it solves; which core entities it uses; what new intelligence it creates; which methodology or rules it relies on; what evidence supports it; whether it is configurable; whether it belongs in TrustLedger or an external system; and how it can be removed or evolved without damaging the core.**

23. What Part 2 Does Not Decide

This capability map does not yet decide the final product scope, commercial packaging, UI navigation, detailed data model, DocTypes, APIs, scoring formulas, AI models, or implementation sequence.

Those decisions belong to later architecture and product reviews after the capability landscape has been examined in sufficient depth.

24. Part 2 Conclusion

TrustLedger should be broad in architectural capacity but selective in product scope.

The stable core should represent the durable reality of stakeholder relationships and project/community interaction. Intelligence should remain modular. Professional methodologies should be pluggable. Client variation should be configurable. Specialist capabilities should be integrated where appropriate. And some adjacent capabilities should deliberately remain outside TrustLedger.

The test is not “Can TrustLedger build this?” The test is “Should TrustLedger own this capability, and can it do so without compromising the architecture's long-term purpose?”

Part 2 therefore establishes the preliminary boundary map for future capability decisions.

Appendix A — Preliminary TrustLedger Capability Landscape

Core operational spine

Organisation → Project / Programme → Place / Context → Stakeholder → Relationship → Engagement → Commitment → Evidence → Decision → Action → Outcome → Learning

Intelligence domains

Stakeholder Intelligence • Relationship Intelligence • Community Intelligence • Engagement & Participation Intelligence • Risk & Early Warning Intelligence • Impact & Social Performance Intelligence • Decision Intelligence

Extensible edge

Professional methodologies • Rules • Models • Indicators • AI assistance • Client configurations • External data • Integrations • Specialist services

Appendix B — Architectural Guardrail

Build the capability when the need is proven. Build the architecture before the capability is needed.

The purpose of this principle is not to predict every future SRM practice. It is to ensure that when a validated new practice emerges, TrustLedger has a place for it without having to dismantle what already works.