

TrustLedger SRM Intelligence Architecture Review

Part 5: Target Intelligence Architecture & Build Blueprint

Working Review Version 0.1 — September 2026

Purpose: Final pre-build architecture decision document. It converts Parts 1–4 into a disciplined target-state blueprint for extending TrustLedger without destabilising its existing product foundation.

1. Purpose and Architectural Decision

Parts 1–4 established the principles, capability boundaries, intelligence model, traceability model and methodology architecture. Part 5 converts those principles into build decisions.

The objective is not to redesign TrustLedger. It is to determine what the existing platform should **keep**, what should be **extended**, what requires **refactoring**, what should be **built**, what belongs through **integration**, what should be **deferred**, and what should explicitly **not be built**.

The result is a controlled path from architecture review into implementation.

2. Target Architectural Proposition

TrustLedger should mature into an **evidence-led Stakeholder Relationship Intelligence platform** for complex projects and programmes.

Its stable operational spine remains:

Organisation → Project/Programme → Place/Context → Stakeholder → Relationship → Engagement → Commitment
→ Evidence → Decision → Action → Outcome → Learning.

Intelligence is not another isolated module. It is the analytical layer operating across this spine.

3. Target Architecture — Layered Model

The target architecture is organised into five broad layers:

Layer 1 — Operational Foundation: organisation/tenancy, projects, places, stakeholders, relationships, engagements, commitments, issues/grievances, evidence, permissions and audit.

Layer 2 — Context & Professional Knowledge: stakeholder operating context, community context, BAU pressure, capacity, competing priorities, accountability, methodologies, frameworks, indicators and professional knowledge.

Layer 3 — Intelligence: signals, interpretation, stakeholder intelligence, relationship intelligence, community intelligence, engagement intelligence, risk/early warning and impact/social-performance intelligence.

Layer 4 — Decision & Action: recommendations, human decisions, actions, commitments, escalation, outcomes and learning.

Layer 5 — Extension & Evolution: integrations, external data, client-specific methodologies, optional capabilities, future professional tools and AI assistance.

The architecture is deliberately layered so that new professional capabilities can be added without rewriting the operational foundation.

4. Build Decision Vocabulary

Every future capability should receive one primary architectural disposition:

KEEP — Existing capability is fundamentally sound and remains part of the stable core.

EXTEND — Existing capability is sound but needs additional fields, relationships, workflows or intelligence without changing its fundamental role.

REFACTOR — Existing implementation creates architectural friction, duplication, false certainty or future extension risk and should be structurally improved before further expansion.

BUILD — Capability is sufficiently justified and absent or materially incomplete; it becomes a controlled new capability.

INTEGRATE — Capability is better provided by an external system/data source while TrustLedger consumes the relevant information.

DEFER — Valuable, but insufficiently proven, operationally dependent or not required for the next maturity step.

DO NOT BUILD — Capability conflicts with TrustLedger's product boundary or would duplicate systems that should remain external.

5. Stable Core — KEEP

The following remain the architectural foundation and should not be casually redesigned:

- Organisation, tenancy, identity and permissions
- Project and programme context
- Place/geographic context
- Stakeholder records
- Engagement records
- Commitments
- Issues/grievances
- Evidence and provenance
- Audit/governance
- Human decision and action records
- Outcome and learning concepts

The core must remain deliberately small. Its job is to represent operational reality, not every possible SRM methodology.

6. Core Extensions — EXTEND

The highest-value extensions are those that make the existing spine more intelligence-ready without turning it into a collection of disconnected modules.

Stakeholder: extend with time-aware operating context such as influence, impact, capacity, commitment, BAU pressure, competing priorities, availability and accountability clarity.

Relationship: strengthen relationship representation, relationship health and traceable relationship evidence.

Engagement: strengthen participation context, burden/load, purpose, attendance/representation where appropriate, and outcome linkage.

Commitment: strengthen ownership, dependency, feasibility, evidence and outcome linkage.

Evidence: strengthen provenance, source, timestamp, context and traceability.

Place/Context: progressively improve geographic and socio-economic context while keeping external data and place packs separable from tenant operational records.

7. Architecture Items Requiring REFACTOR

Refactoring should be selective and evidence-led. The current review identifies areas where implementation patterns could constrain the target architecture:

- Local/browser-only authored-report persistence where durable organisational reporting is required.
- Geo architecture where local place packs and live Cloud retrieval currently coexist without a complete first-class operational model.
- Intelligence paths that can fall back to generic/demo data when identifiers or context are incomplete.
- Intelligence implementations where domain logic is tightly coupled to UI or mock-data assumptions.
- Relationship/influence representation that exists conceptually but is not yet a sufficiently general relationship model.

Refactoring should not be undertaken merely to make code cleaner. Each refactor must have a specific architectural purpose and acceptance criteria.

8. Priority BUILD Capabilities

The target build sequence should prioritise the intelligence infrastructure rather than a large number of visible features.

Build Priority 1 — Common Intelligence Object / Traceability Pattern: establish the conceptual and implementation contract for fact, observation, signal, interpretation, intelligence, recommendation, decision, action, outcome and learning.

Build Priority 2 — Contextual Intelligence: make stakeholder operating context and relevant project/community context usable as intelligence inputs.

Build Priority 3 — Relationship Intelligence: mature relationship representation, relationship health and explainable relationship signals.

Build Priority 4 — Signal and Early-Warning Architecture: create explainable signals from changes in commitments, grievances, engagement patterns, stakeholder context and other evidence.

Build Priority 5 — Decision Intelligence: connect intelligence to recommendations, human decisions, actions and outcomes without automating accountability.

These are architecture priorities, not a permission to build every item immediately.

9. INTEGRATE Rather Than Rebuild

TrustLedger should integrate with systems that are authoritative or specialised outside SRM.

Potential integration domains include:

- GIS and authoritative geographic datasets
- CRM/customer systems where stakeholder identity must be synchronised
- ERP/project controls for programme or delivery information
- External socio-economic/statistical datasets
- Document/evidence repositories where appropriate
- AI services where governed and privacy-safe
- Client enterprise systems

The rule is: **TrustLedger should own the relationship intelligence created from relevant information, not necessarily the source system that originated every data point.**

10. DEFER

The following should remain deferred until there is stronger evidence, operational demand or architectural readiness:

- Advanced predictive forecasting
- Full stakeholder network visualisation/graph analytics
- Social-media monitoring as a major product subsystem
- A broad methodology marketplace
- Highly automated AI decisioning
- Comprehensive ESG management
- Large-scale programme-management functionality
- A fully general-purpose community development management system
- Complex statistical modelling before reliable longitudinal data exists

Deferral is not rejection. It protects the architecture from premature complexity.

11. DO NOT BUILD — Product Boundary

TrustLedger should not attempt to become:

- A generic CRM
- An ERP
- A GIS platform
- A project-management replacement
- A generic survey platform
- A generic AI platform
- A complete ESG suite
- A social-media intelligence platform
- A community-development administration system
- A replacement for professional judgement or accountable governance

TrustLedger may integrate with these domains and derive intelligence from them, but its differentiated role is the relationship/social intelligence layer.

12. Methodology Architecture — Target State

Methodologies remain pluggable analytical lenses.

A methodology should enter the intelligence pipeline through a controlled contract:

Methodology identity → version → scope → required inputs → analytical logic → outputs → evidence requirements → limitations → governance.

This allows future approaches—stakeholder salience, power-interest, influence networks, participation, conflict, social risk, social performance, Social Licence to Build, IKS/local knowledge and client-specific frameworks—to coexist without hard-coding one doctrine into the core.

13. AI Position in the Target Architecture

AI is an **assistive capability**, not the source of truth.

Target flow:

Evidence/Data → Context → Rules/Methodologies/Models → Intelligence → AI Assistance → Human Review → Decision → Action → Outcome.

AI may assist with extraction, summarisation, classification, drafting, pattern identification and recommendation generation where authorised.

AI must not silently create authoritative facts, bypass human governance, expose tenant data, or convert uncertain interpretation into false certainty.

14. Genuine Data vs Demo/Mock Boundary

The target architecture requires a strict separation between:

Authoritative tenant data — operational customer data.

Reference/external data — clearly identified contextual datasets.

Methodology outputs — analytical results with provenance.

Demo/sample data — explicitly non-customer content.

No intelligence feature should silently substitute demo/sample records for missing customer evidence. Empty authoritative data must remain empty. Any fallback must be explicit, bounded and visible to the system rather than masquerading as customer reality.

15. Intelligence Lifecycle — Target

The target lifecycle is:

DATA → CONTEXT → SIGNAL → INTERPRETATION → INTELLIGENCE → RECOMMENDATION → HUMAN DECISION → ACTION → OUTCOME → LEARNING.

Every major intelligence capability should be traceable through this chain.

A signal is not a conclusion. An interpretation is not a fact. A recommendation is not a decision. A decision is not an outcome.

This separation is one of the central trust mechanisms of TrustLedger.

16. Build Order

The recommended build order is:

Stage 1 — Protect the core. Confirm tenancy, identity, operational entities, evidence and audit boundaries.

Stage 2 — Establish common intelligence infrastructure. Introduce shared contracts/patterns rather than separate intelligence silos.

Stage 3 — Extend context. Add stakeholder operating context and contextual inputs.

Stage 4 — Mature relationships. Build relationship intelligence on evidence already captured.

Stage 5 — Build explainable signals. Use rules, methodologies and evidence to detect meaningful change.

Stage 6 — Connect decisions and outcomes. Close the loop between intelligence, action and learning.

Stage 7 — Add specialised intelligence. Introduce community, participation, risk and social-performance capabilities where justified.

Stage 8 — Integrate advanced AI/external intelligence. Only after the underlying evidence and governance foundations are reliable.

17. Capability Admission Gate

Before any future capability enters development, ask:

1. Is it fundamentally relevant to Stakeholder Relationship Management? 2. Does it strengthen TrustLedger's trust-building purpose? 3. Does it connect to the operational spine? 4. Can it be added as an extension, methodology,

intelligence capability or integration rather than destabilising the core? 5. Is there a real user/problem need? 6. Is the evidence sufficient? 7. Is there a clear owner and human governance path? 8. Can its outputs be explained and traced? 9. Can it coexist with alternative methodologies? 10. What should TrustLedger explicitly *not* build to support it?

A “yes” does not automatically authorise development. It establishes architectural eligibility.

18. Definition of Architectural Success

The architecture succeeds if TrustLedger can add a genuinely new SRM capability—arising from professional practice, client demand, new evidence, regulation, technology or emerging methodology—without requiring redesign of already-stable core capabilities.

A new capability should be able to attach to the existing spine, consume common evidence/context, apply its own methodology or intelligence logic, produce traceable outputs, and feed human decisions and outcomes.

19. Final Target-State Pattern

The target pattern is:

Stable Operational Core ↓ Common Evidence & Context ↓ Signals / Interpretations ↓ Modular Intelligence Domains ↓ Methodologies / Rules / Models ↓ Intelligence Synthesis ↓ Recommendations ↓ Human Decision & Governance ↓ Action / Commitment ↓ Outcome / Learning

External systems, external data and AI surround this architecture as controlled extensions—not as replacements for its core.

20. What This Document Authorises

This document establishes the **target architecture and build decision framework**.

It does **not** by itself authorise implementation of every capability named here.

Before coding begins, each implementation packet should identify:

- architectural decision (KEEP / EXTEND / REFACTOR / BUILD / INTEGRATE / DEFER / DO NOT BUILD)
- scope
- affected existing capability
- data ownership
- evidence/provenance requirements
- methodology implications
- tenancy/security implications
- human governance
- acceptance criteria
- verification method
- rollback or containment strategy where applicable

This preserves the DESIGN → ASSIGN → EXECUTE → HANDOFF → VERIFY → CLOSE → EMPTY → NEXT TASK discipline.

21. Final Architecture Decision

TrustLedger now has sufficient strategic architecture to move from open-ended architecture exploration into controlled implementation planning.

The guiding rule is:

Build the capability when the need is proven. Build the architecture before the capability is needed.

TrustLedger should therefore be **comprehensive in architectural capacity, selective in product scope, evidence-led in intelligence, methodology-agnostic, human-governed and extensible at the edge while keeping its operational core stable.**

Appendix A — Target Capability Disposition Matrix

Capability area	Disposition	Target treatment
Organisation / tenancy / identity	KEEP	Protect as stable foundation
Projects / programmes	KEEP + EXTEND	Preserve core; improve context linkage
Place / geographic context	EXTEND + INTEGRATE	Strengthen operational context; consume authoritative external data
Stakeholders	EXTEND	Add time-aware operating context
Relationships	BUILD + EXTEND	Mature relationship intelligence
Engagements	EXTEND	Strengthen participation and outcome linkage
Commitments	EXTEND	Add feasibility, evidence and dependency intelligence
Issues / grievances	KEEP + EXTEND	Use as evidence and early-warning source
Evidence / provenance	EXTEND	Make traceability a shared intelligence foundation
Reporting	REFACTOR + EXTEND	Move toward durable organisational reporting
Intelligence object / traceability	BUILD	Shared cross-domain intelligence contract
Stakeholder operating context	BUILD	Contextual intelligence foundation
Signal / early warning	BUILD	Explainable evidence-based signals
Decision intelligence	BUILD	Recommendations linked to human decisions and outcomes
Methodology architecture	BUILD	Pluggable methodology contract and governance
Advanced predictive AI	DEFER	Require longitudinal evidence and governance
External systems / datasets	INTEGRATE	Use authoritative specialised systems
Generic CRM / ERP / GIS replacement	DO NOT BUILD	Remain outside product boundary

Appendix B — Pre-Build Gate

Architecture review sequence complete: Blueprint → Intelligence Model → Capability Map → Traceability Model → Methodology & Framework Architecture → Target Intelligence Architecture & Build Blueprint.

Recommended next step: stop architecture-document production and begin a controlled implementation planning phase. The first implementation task should be selected from the target architecture and assigned through the established .ai/TASK.md gate; no broad “build everything” instruction should be issued.